

In [1]:

```
import tensorflow as tf
```

In [2]:

```
tf.__version__
```

Out[2]:

```
'2.0.0'
```

In [35]:

```
from tensorflow.keras import layers, optimizers, metrics, datasets, Sequential, models
```

In [4]:

```
import os
```

In [5]:

```
import matplotlib.pyplot as plt  
%matplotlib inline
```

In [9]:

```
(x_train, y_train), (x_test, y_test) = datasets.imdb.load_data(num_words=10000)
```

In [10]:

```
x_train.shape, y_train.shape, x_test.shape, y_test.shape
```

Out[10]:

```
((25000,), (25000,), (25000,), (25000,))
```

In [14]:

```
x_train[0][:10]
```

Out[14]:

```
[1, 14, 22, 16, 43, 530, 973, 1622, 1385, 65]
```

In [16]:

```
y_train[0]
```

Out[16]:

```
1
```

In [17]:

```
max([max(sequence) for sequence in x_train])
```

Out[17]:

9999

数字和单词映射表，索引减3，因为0, 1, 2为padding、start of sequence、unknown保留的索引

In [19]:

```
word_index = datasets.imdb.get_word_index()
reverse_word_index = dict([(value, key) for (key, value) in word_index.items()])
decoded_review = ' '.join([reverse_word_index.get(i - 3, '?') for i in x_train[0]])
```

In [20]:

```
decoded_review
```

Out[20]:

```
"? this film was just brilliant casting location scenery story direction everyone's really suited the part they played and you could just imagine being there robert ? is an amazing actor and now the same being director ? father came from the same scottish island as myself so i loved the fact there was a real connection with this film the witty remarks throughout the film were great it was just brilliant so much that i bought the film as soon as it was released for ? and would recommend it to everyone to watch and the fly fishing was amazing really cried at the end it was so sad and you know what they say if you cry at a film it must have been good and this definitely was also ? to the two little boy's that played the ? of norman and paul they were just brilliant children are often left out of the ? list i think because the stars that play them all grown up are such a big profile for the whole film but these children are amazing and should be praised for what they have done don't you think the whole story was so lovely because it was true and was someone's life after all that was shared with us all"
```

In [21]:

```
import numpy as np
```

向量化

In [22]:

```
def vectorize_sequences(sequences, dimension=10000):
    results = np.zeros((len(sequences), dimension))
    for i, sequence in enumerate(sequences):
        results[i, sequence] = 1
    return results
```

In [26]:

```
x_train = vectorize_sequences(x_train)
x_test = vectorize_sequences(x_test)
```

In [29]:

```
x_train.shape
```

Out[29]:

```
(25000, 10000)
```

In [30]:

```
x_train[0].shape
```

Out[30]:

```
(10000,)
```

In [31]:

```
x_train[0]
```

Out[31]:

```
array([0., 1., 1., ..., 0., 0., 0.])
```

In [32]:

```
y_train = np.asarray(y_train).astype('float32')
y_test = np.asarray(y_test).astype('float32')
```

In [33]:

```
y_train[0]
```

Out[33]:

```
1.0
```

In [36]:

```
model = models.Sequential()
model.add(layers.Dense(16, activation='relu', input_shape=(10000, )))
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

In [37]:

```
model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['accuracy'])
```

In [42]:

```
# model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
```

In [38]:

```
# model.compile(optimizer=optimizers.RMSprop(lr=0.001), loss='binary_crossentropy', metrics
```

In [40]:

```
# model.compile(optimizer=optimizers.RMSprop(lr=0.001), loss=losses.binary_crossentropy, me
```

留出验证集

In [41]:

```
x_val = x_train[:10000]  
x_train = x_train[10000:]  
  
y_val = y_train[:10000]  
y_train = y_train[10000:]
```

In [43]:

```
history = model.fit(x_train, y_train, epochs=20, batch_size=512, validation_data=(x_val, y_val))
```

Train on 15000 samples, validate on 10000 samples

Epoch 1/20

15000/15000 [=====] - 2s 122us/sample - loss: 0.6140 - accuracy: 0.5672 - val_loss: 0.5530 - val_accuracy: 0.7330

Epoch 2/20

15000/15000 [=====] - 1s 54us/sample - loss: 0.5005 - accuracy: 0.8133 - val_loss: 0.4925 - val_accuracy: 0.8268

Epoch 3/20

15000/15000 [=====] - 1s 54us/sample - loss: 0.4313 - accuracy: 0.8904 - val_loss: 0.4632 - val_accuracy: 0.8236

Epoch 4/20

15000/15000 [=====] - 1s 53us/sample - loss: 0.3568 - accuracy: 0.9241 - val_loss: 0.3833 - val_accuracy: 0.8882

Epoch 5/20

15000/15000 [=====] - 1s 54us/sample - loss: 0.2721 - accuracy: 0.9447 - val_loss: 0.3422 - val_accuracy: 0.8810

Epoch 6/20

15000/15000 [=====] - 1s 53us/sample - loss: 0.2024 - accuracy: 0.9565 - val_loss: 0.3038 - val_accuracy: 0.8859

Epoch 7/20

15000/15000 [=====] - 1s 53us/sample - loss: 0.1558 - accuracy: 0.9657 - val_loss: 0.3056 - val_accuracy: 0.8831

Epoch 8/20

15000/15000 [=====] - 1s 53us/sample - loss: 0.1217 - accuracy: 0.9748 - val_loss: 0.3037 - val_accuracy: 0.8823

Epoch 9/20

15000/15000 [=====] - 1s 53us/sample - loss: 0.0974 - accuracy: 0.9793 - val_loss: 0.3241 - val_accuracy: 0.8804

Epoch 10/20

15000/15000 [=====] - 1s 53us/sample - loss: 0.0817 - accuracy: 0.9819 - val_loss: 0.3411 - val_accuracy: 0.8779

Epoch 11/20

15000/15000 [=====] - 1s 53us/sample - loss: 0.0644 - accuracy: 0.9879 - val_loss: 0.3544 - val_accuracy: 0.8785

Epoch 12/20

15000/15000 [=====] - 1s 54us/sample - loss: 0.0538 - accuracy: 0.9895 - val_loss: 0.4190 - val_accuracy: 0.8711

Epoch 13/20

15000/15000 [=====] - 1s 54us/sample - loss: 0.0452 - accuracy: 0.9921 - val_loss: 0.3974 - val_accuracy: 0.8774

Epoch 14/20

15000/15000 [=====] - 1s 54us/sample - loss: 0.0363 - accuracy: 0.9940 - val_loss: 0.4267 - val_accuracy: 0.8756

Epoch 15/20

15000/15000 [=====] - 1s 53us/sample - loss: 0.0302 - accuracy: 0.9954 - val_loss: 0.4549 - val_accuracy: 0.8735

Epoch 16/20

15000/15000 [=====] - 1s 54us/sample - loss: 0.0240 - accuracy: 0.9963 - val_loss: 0.4864 - val_accuracy: 0.8731

Epoch 17/20

15000/15000 [=====] - 1s 54us/sample - loss: 0.0188 - accuracy: 0.9974 - val_loss: 0.5344 - val_accuracy: 0.8687

Epoch 18/20

15000/15000 [=====] - 1s 54us/sample - loss: 0.0146 - accuracy: 0.9981 - val_loss: 0.5372 - val_accuracy: 0.8706

Epoch 19/20

```
15000/15000 [=====] - 1s 55us/sample - loss: 0.0  
122 - accuracy: 0.9983 - val_loss: 0.6691 - val_accuracy: 0.8610  
Epoch 20/20  
15000/15000 [=====] - 1s 55us/sample - loss: 0.0  
100 - accuracy: 0.9985 - val_loss: 0.6051 - val_accuracy: 0.8699
```

In [44]:

```
history_dict = history.history  
history_dict.keys()
```


Out[44]:

```
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

绘图

In [47]:

```
loss = history_dict['loss']  
val_loss = history_dict['val_loss']  
  
epochs = range(1, len(loss) + 1)  
  
plt.plot(epochs, loss, 'bo', label='Training loss')  
plt.plot(epochs, val_loss, 'b', label='Validation loss')  
plt.title('Training and validation loss')  
plt.xlabel('Epochs')  
plt.ylabel('Loss')  
plt.legend()  
  
plt.show()
```




In [48]:

```
# plt.clf()      # 清除图像
acc = history_dict['accuracy']
val_acc = history_dict['val_accuracy']

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```



```
loss, acc = model.evaluate(x_test, y_test)
```

In [50]:

loss, acc

 $(0.6686725060486793, 0.85684)$

```
model.predict(x_test)
```

```
array([[0.0075345 ],
       [0.9999391 ],
       [0.43583608],
       ...,
       [0.00217295],
       [0.00470984],
       [0.6878468 ]], dtype=float32)
```

early stopping, 只训练4个周期

In []:

```
model = models.Sequential()
model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))

model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['accuracy'])

model.fit(x_train, y_train, epochs=4, batch_size=512)
loss, acc = model.evaluate(x_test, y_test)
loss, acc
```